

## Problem 3.24

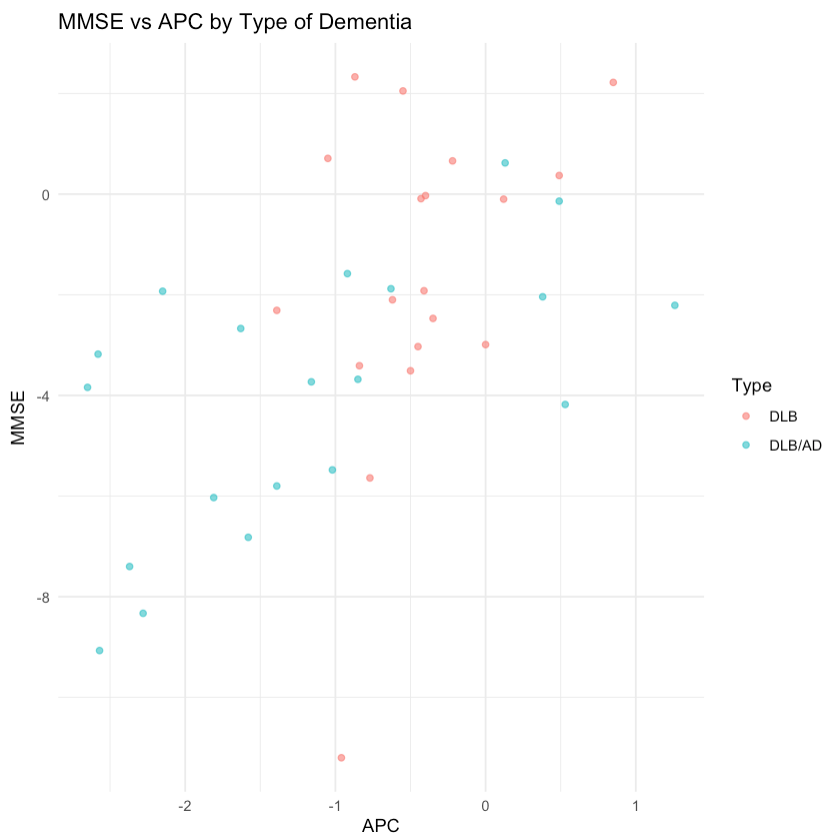
a. Make a scatterplot of  $Y = \text{MMSE}$  versus  $X = \text{APC}$  and with *Type* as a grouping variable (i.e., use different colors or different plotting symbols for the two levels of *Type*)

note: add a part d (look at assignment 5 instructions)

```
In [1]: lewy_data <- read.csv("../data/LewyBody2Groups.csv")

library(ggplot2)

ggplot(lewy_data, aes(x = APC, y = MMSE, color = Type)) +
  geom_point(alpha = 0.6) +
  labs(
    title = "MMSE vs APC by Type of Dementia",
    x = "APC",
    y = "MMSE"
  ) +
  theme_minimal()
```



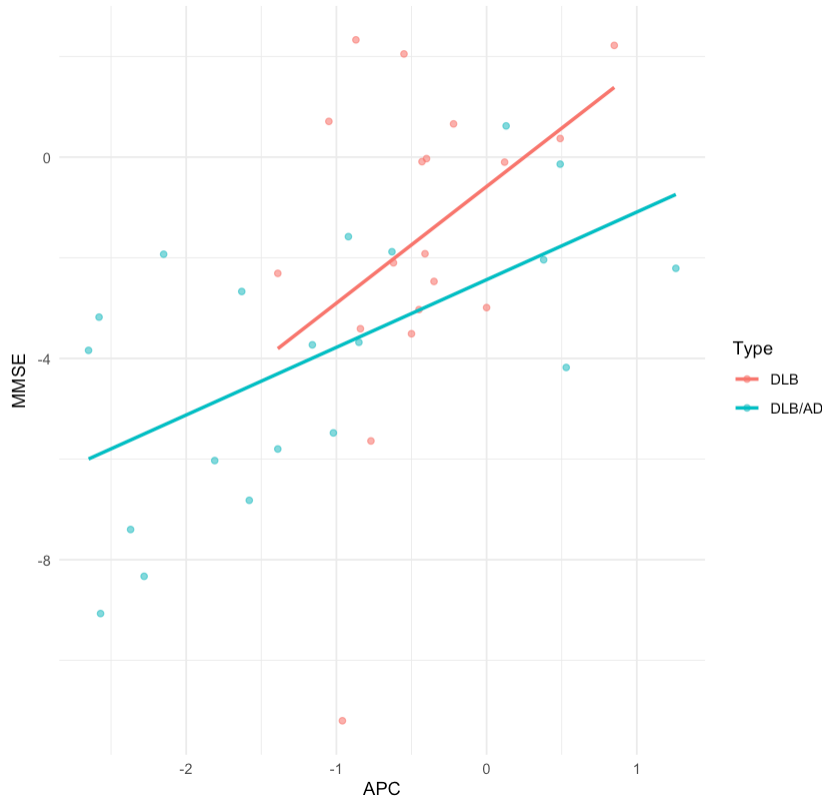
There appears to be a positive linear relationship between the two variables with both types, with DLB having a sharper slope. We can see this if we map the data to a linear model for each type and overlay it on our scatterplot:

```
In [3]: ggplot(lewy_data, aes(x = APC, y = MMSE, color = Type)) +
```

```
geom_point(alpha = 0.6) +
geom_smooth(method = "lm", se = FALSE) +
labs(
  title = "MMSE vs APC by Type of Dementia",
  x = "APC",
  y = "MMSE"
) +
theme_minimal()
```

`geom\_smooth()` using formula = 'y ~ x'

MMSE vs APC by Type of Dementia



b. Fit the regression of MMSE on APC and test whether there is a linear association between the two variables

We can use a linear model to fit the regression of MMSE on APC, and then apply a test of correlation to both variables to assess the strength of the relationship:

```
In [4]: model <- lm(MMSE ~ APC, data = lewy_data)
summary(model)

cor.test(lewy_data$MMSE, lewy_data$APC)
```

```

Call:
lm(formula = MMSE ~ APC, data = lewy_data)

Residuals:
    Min       1Q   Median       3Q      Max
-8.1022 -1.7043  0.2174  1.9484  5.2706

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -1.4214     0.5528  -2.572  0.014277 *
APC           1.7462     0.4401   3.968  0.000321 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.664 on 37 degrees of freedom
Multiple R-squared:  0.2985,    Adjusted R-squared:  0.2795
F-statistic: 15.74 on 1 and 37 DF,  p-value: 0.0003208
        Pearson's product-moment correlation

data:  lewy_data$MMSE and lewy_data$APC
t = 3.9675, df = 37, p-value = 0.0003208
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.2788694 0.7351198
sample estimates:
      cor
0.5463171

```

Overall, we see that there is a significant positive linear association between MMSE and APC, and it's a moderately strong one (given by our 0.546 correlation coefficient).

c. Fit a model that produces parallel regression lines for the two levels of *Type*. Write down the fitted prediction equation for each level of *Type*.

```
In [7]: parallel_model <- lm(MMSE ~ APC + Type, data = lewy_data)
summary(parallel_model)
```

```

Call:
lm(formula = MMSE ~ APC + Type, data = lewy_data)

Residuals:
    Min       1Q   Median       3Q      Max
-8.8153 -1.6382 -0.1469  1.9103  4.5796

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -0.9433     0.6358  -1.484  0.14662
APC           1.5015     0.4650   3.229  0.00265 **
TypeDLB/AD   -1.3135     0.9017  -1.457  0.15385
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.624 on 36 degrees of freedom
Multiple R-squared:  0.3375,    Adjusted R-squared:  0.3007
F-statistic:  9.17 on 2 and 36 DF,  p-value: 0.0006042

```

Using the output above, we see that the slope is 1.5015 for APC, with an intercept of -0.9433. The DLB/AD type's intercept is further decreased by 1.3135, with a final intercept of -2.2568. Putting it all together, we get the following prediction equations:

For DLB:

$$MMSE = -0.9433 + 1.5015(APC)$$

For DLB/AD:

$$MMSE = -2.2568 + 1.5015(APC)$$

One important observation is that the APC coefficient (1.5015) is significant, but the difference between the two levels (DLB vs DLB/AD) is not (.15385)

d. Create a model with interaction of the IType variable with the variable APC so that the regression lines can have different slopes and intercepts.

```
In [8]: interaction_model <- lm(MMSE ~ APC + Type + APC:Type, data = lewy_data)
summary(interaction_model)

ggplot(lewy_data, aes(x = APC, y = MMSE, color = Type)) +
  geom_point(alpha = 0.6) +
  geom_smooth(method = "lm", se = FALSE) +
  labs(
    title = "MMSE vs APC by Type with Different Slopes",
    x = "APC",
    y = "MMSE"
  ) +
  theme_minimal()
```

Call:

```
lm(formula = MMSE ~ APC + Type + APC:Type, data = lewy_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-8.3905	-1.5841	-0.1014	1.6959	4.9309

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.5846	0.7927	-0.738	0.4657
APC	2.3176	1.1640	1.991	0.0543 .
TypeDLB/AD	-1.8513	1.1471	-1.614	0.1155
APC:TypeDLB/AD	-0.9732	1.2712	-0.766	0.4490

----

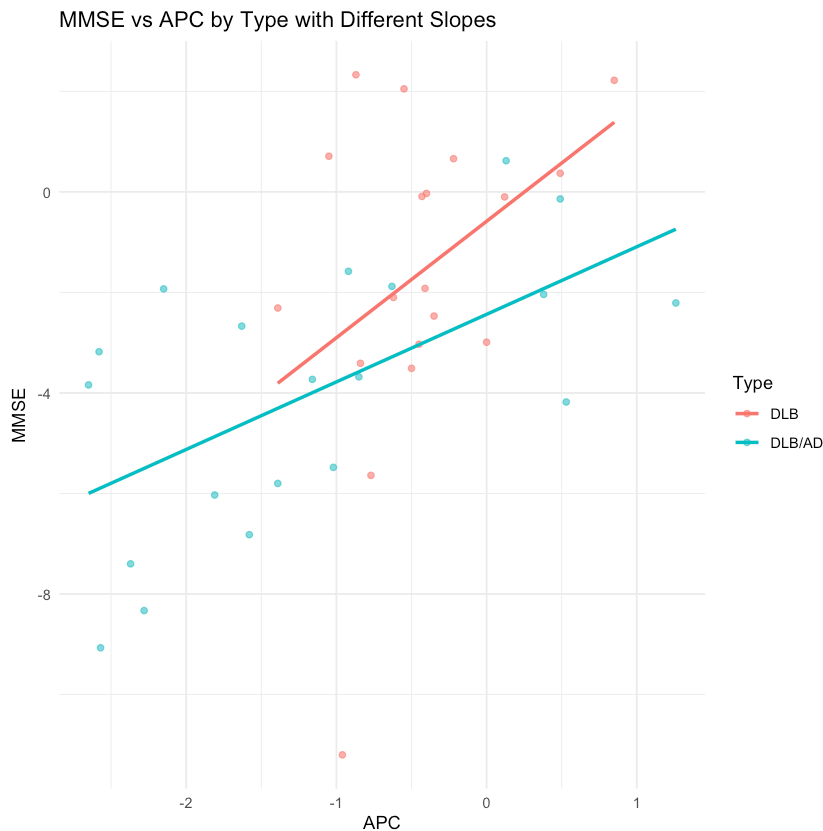
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.64 on 35 degrees of freedom

Multiple R-squared: 0.3484, Adjusted R-squared: 0.2926

F-statistic: 6.239 on 3 and 35 DF, p-value: 0.001656

`geom\_smooth()` using formula = 'y ~ x'



## Problem 3.38

a. Run a quadratic model using Depth

```
In [10]: diamonds <- read.csv("../data/Diamonds.csv")
diamonds$Depth_sq <- diamonds$Depth^2

quad_model <- lm(TotalPrice ~ Depth + Depth_sq, data = diamonds)
summary(quad_model)

ggplot(diamonds, aes(x = Depth, y = TotalPrice)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = "lm", formula = y ~ poly(x, 2), se = FALSE) +
  labs(
    title = "Diamond Price vs Depth (Quadratic Fit)",
    x = "Depth",
    y = "Total Price"
  ) +
  theme_minimal()
```

Call:

```
lm(formula = TotalPrice ~ Depth + Depth_sq, data = diamonds)
```

Residuals:

Min	1Q	Median	3Q	Max
-9323	-4251	-2676	2134	45513

Coefficients:

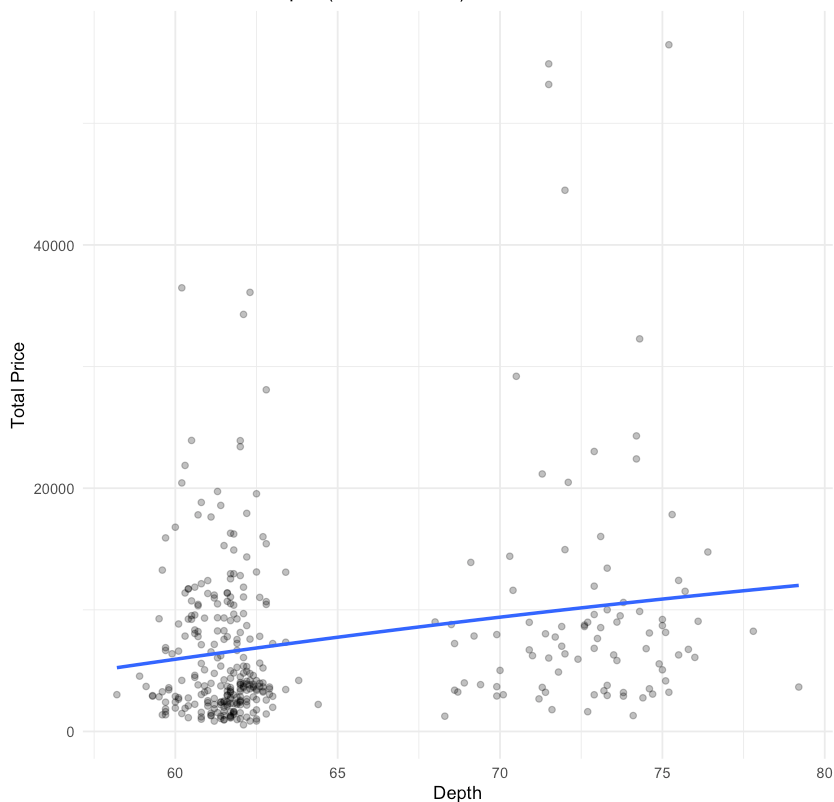
	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-28406.783	112211.790	-0.253	0.800
Depth	766.369	3353.222	0.229	0.819
Depth_sq	-3.233	24.869	-0.130	0.897

Residual standard error: 7616 on 348 degrees of freedom

Multiple R-squared: 0.04748, Adjusted R-squared: 0.042

F-statistic: 8.673 on 2 and 348 DF, p-value: 0.0002111

Diamond Price vs Depth (Quadratic Fit)



$$R^2 = 0.04748$$

$$\text{Adjusted } R^2 = 0.042$$

Neither depth ( $p=.819$ ) nor  $depth^2$  ( $p=.897$ ) are statistically significant.

b. A two-predictor model using Carat and Depth

```
In [13]: two_pred_model <- lm(TotalPrice ~ Carat + Depth, data = diamonds)
summary(two_pred_model)

ggplot(diamonds, aes(x = Depth, y = TotalPrice)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = "lm", se = FALSE) +
```

```
labs(
  title = "Diamond Price vs Depth (controlling for Carat)",
  x = "Depth",
  y = "Total Price"
) +
theme_minimal()
```

Call:

```
lm(formula = TotalPrice ~ Carat + Depth, data = diamonds)
```

Residuals:

Min	1Q	Median	3Q	Max
-9234.7	-1223.7	-274.3	1161.0	16368.6

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	1059.24	1918.36	0.552	0.581
Carat	15087.01	320.96	47.006	< 2e-16 ***
Depth	-134.94	30.92	-4.364	1.68e-05 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

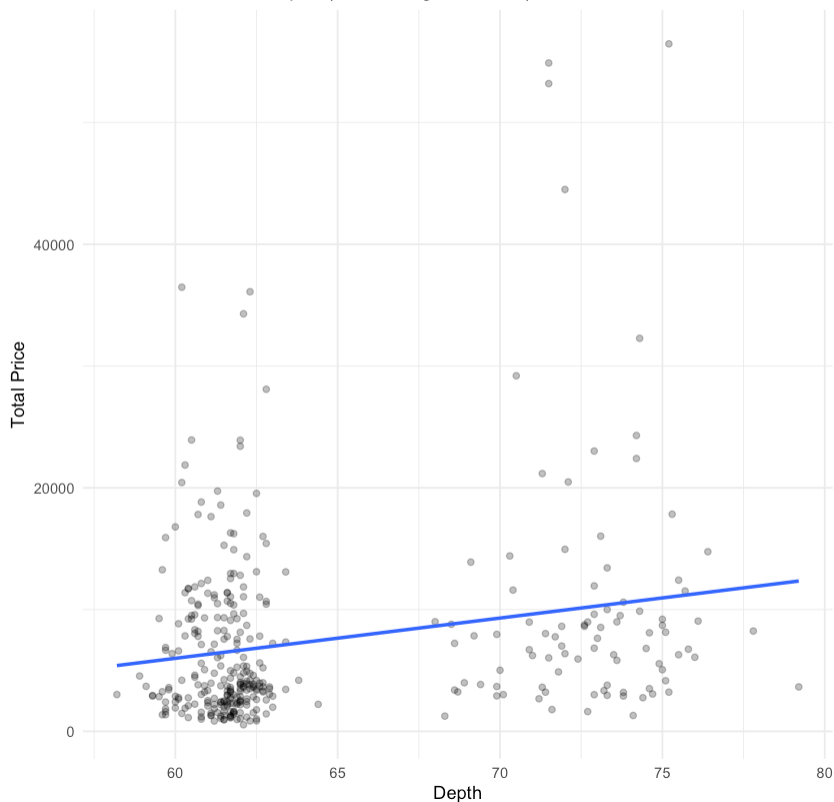
Residual standard error: 2809 on 348 degrees of freedom

Multiple R-squared: 0.8704, Adjusted R-squared: 0.8696

F-statistic: 1168 on 2 and 348 DF, p-value: < 2.2e-16

`geom\_smooth()` using formula = 'y ~ x'

Diamond Price vs Depth (controlling for Carat)



$$R^2 = 0.8704$$

$$\text{Adjusted } R^2 = 0.8696$$

This is statistically significant model, with an F-statistic p-value  $< 2.2\text{e-}16$ . Additionally, unlike with the quadratic model, our coefficients are statistically significant:

Carat:  $< 2\text{e-}16$

Depth:  $< 1.68\text{e-}05$

This model performs considerably better than the quadratic model that used just depth, ( $R^2$  increases from .046 to .87), and we note that carat has a much stronger effect on price than depth.

c. A three-predictor model that adds interaction for *Carat* and *Depth*

```
In [14]: interaction_model <- lm(TotalPrice ~ Carat + Depth + Carat:Depth, data = dia
summary(interaction_model)

diamonds$CaratGroup <- cut(diamonds$Carat,
                           breaks = c(0, 0.5, 1, 1.5, 2, max(diamonds$Carat)),
                           labels = c("0-0.5", "0.5-1", "1-1.5", "1.5-2", "2+"))

ggplot(diamonds, aes(x = Depth, y = TotalPrice, color = CaratGroup)) +
  geom_point(alpha = 0.3) +
  geom_smooth(method = "lm", se = FALSE) +
  labs(
    title = "Diamond Price vs Depth by Carat Group",
    x = "Depth",
    y = "Total Price",
    color = "Carat Range"
  ) +
  theme_minimal()
```

Call:

```
lm(formula = TotalPrice ~ Carat + Depth + Carat:Depth, data = diamonds)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-8254.4	-1311.5	-157.2	1131.8	14513.9

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	31171.41	4219.58	7.387	1.13e-12 ***
Carat	-11827.73	3436.47	-3.442	0.000648 ***
Depth	-598.18	65.47	-9.137	< 2e-16 ***
Carat:Depth	408.45	51.96	7.861	4.84e-14 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2592 on 347 degrees of freedom

Multiple R-squared: 0.89, Adjusted R-squared: 0.889

F-statistic: 935.7 on 3 and 347 DF, p-value:  $< 2.2\text{e-}16$

`geom\_smooth()` using formula = 'y ~ x'





$$R^2 = 0.89$$

$$\text{Adjusted } R^2 = 0.889$$

This is statistically significant model, with an F-statistic p-value  $< 2.2e-16$ . Additionally, unlike with the quadratic model and like with our two-predictor model, all of our coefficients are statistically significant:

Carat: .000648

Depth:  $< 2e-16$

Carat:Depth:  $4.84e-14$

The interaction term between Carat and Depth is \$408 per unit of both carat and depth, which suggests that depth becomes more important as carat increases. This model also performs slightly better than the two predictor model (89% of the variation explained, vs 87%).

d. A complete second-order model using Carat and Depth

```
In [15]: diamonds$Carat_sq <- diamonds$Carat^2
diamonds$Depth_sq <- diamonds$Depth^2

complete_model <- lm(TotalPrice ~ Carat + Depth + Carat:Depth + Carat_sq + D
summary(complete_model)

ggplot(diamonds, aes(x = Depth, y = TotalPrice, color = CaratGroup)) +
```

```

geom_point(alpha = 0.3) +
geom_smooth(method = "lm", formula = y ~ poly(x, 2), se = FALSE) +
facet_wrap(~CaratGroup, scales = "free_y") +
labs(
  title = "Diamond Price vs Depth by Carat Group (Second Order)",
  x = "Depth",
  y = "Total Price"
) +
theme_minimal()

```

Call:

```
lm(formula = TotalPrice ~ Carat + Depth + Carat:Depth + Carat_sq +
    Depth_sq, data = diamonds)
```

Residuals:

Min	1Q	Median	3Q	Max
-12196.1	-652.7	-38.5	485.7	10582.2

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	24338.820	30297.912	0.803	0.4223
Carat	7573.620	3040.787	2.491	0.0132 *
Depth	-728.700	904.439	-0.806	0.4210
Carat_sq	4761.592	330.246	14.418	<2e-16 ***
Depth_sq	5.276	6.727	0.784	0.4333
Carat:Depth	-83.891	53.530	-1.567	0.1180

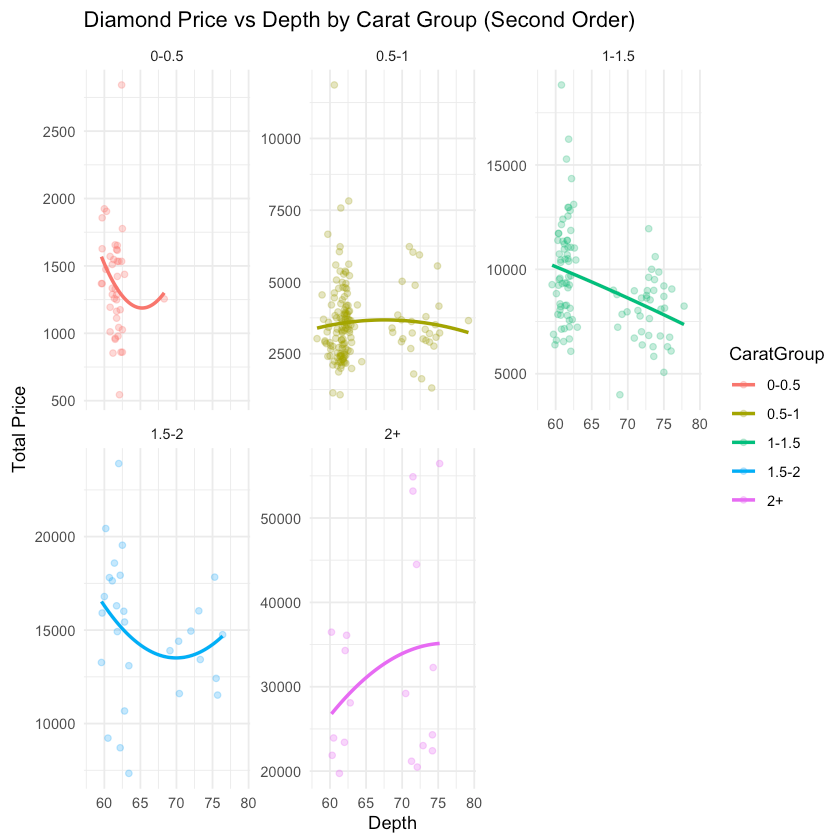
---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2053 on 345 degrees of freedom

Multiple R-squared: 0.9313, Adjusted R-squared: 0.9304

F-statistic: 936.1 on 5 and 345 DF, p-value: < 2.2e-16



$$R^2 = 0.913$$

$$\text{Adjusted } R^2 = 0.9304$$

Again, this is a significant model overall ( $F=936.1$ ,  $p < 2.2e-16$ )

Carat: 0.0132 (significant)

Depth: .4210 (not significant)

$\text{Carat}^2 = < 2e-16$  (significant)

$\text{Depth}^2 = .4333$  (not significant)

Carat x Depth = .118 (not significant)

This model performs the best out of all the models so far (explaining 93% of Total Price). It also suggests that depth is not nearly as important as other models have shown, and the interaction with carat is no longer significant when controlling for squared terms. I would recommend this model for predicting the total price of diamonds, as it provides the best fit while only increasing complexity modestly.

## Problem 3.39

a. Using the model you chose, produce one or more graphs to examine the conditions for homoscedasticity and normality of its residuals. Do these standard regression

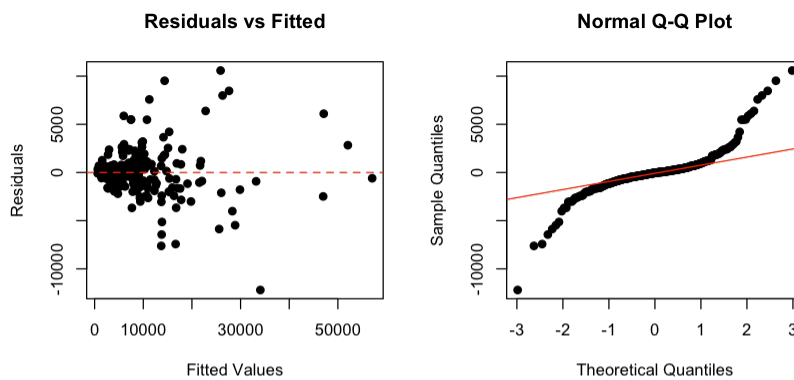
conditions appear to be reasonable for your model?

```
In [16]: residuals <- resid(complete_model)
         fitted_values <- fitted(complete_model)

         par(mfrow=c(2,2))

         plot(fitted_values, residuals,
              xlab="Fitted Values",
              ylab="Residuals",
              main="Residuals vs Fitted",
              pch=19)
         abline(h=0, col="red", lty=2)

         qqnorm(residuals, pch=19)
         qqline(residuals, col="red")
```



Our residuals vs fitted plot shows a fanning pattern moving right, indicating larger residuals for higher fitted values, which suggests some violation of the constant variance assumption. Our QQ plot has relatively good alignment toward the middle, but there are deviations in the tails, particularly at the higher end, which suggests a slight departure from normality.

b. Transform the response variable to be *logPrice* as the natural log of the *TotalPrice*. Is your "best" choice of models still a reasonable choice for predicting *logPrice*? If not, make adjustments to add or delete terms, keeping with the options offered within a complete second-order model.

```
In [19]: diamonds$logPrice <- log(diamonds$TotalPrice)
```

```

log_complete_model <- lm(logPrice ~ Carat + Depth + Carat:Depth + Carat_sq +
summary(log_complete_model)

residuals <- resid(log_complete_model)
fitted_values <- fitted(log_complete_model)

par(mfrow=c(2,2))

plot(fitted_values, residuals,
      xlab="Fitted Values",
      ylab="Residuals",
      main="Residuals vs Fitted (Log Model)",
      pch=19)
abline(h=0, col="red", lty=2)

qqnorm(residuals, pch=19)
qqline(residuals, col="red")

# simpler model without the insignificant terms
log_reduced_model <- lm(logPrice ~ Carat + Carat_sq, data = diamonds)
summary(log_reduced_model)

```

Call:

```
lm(formula = logPrice ~ Carat + Depth + Carat:Depth + Carat_sq +
    Depth_sq, data = diamonds)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.85021	-0.13209	0.01441	0.13613	0.79710

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	13.5049624	3.4020467	3.970	8.76e-05	***
Carat	2.5863485	0.3414393	7.575	3.33e-13	***
Depth	-0.2027689	0.1015563	-1.997	0.0467	*
Carat_sq	-0.5714071	0.0370821	-15.409	< 2e-16	***
Depth_sq	0.0013384	0.0007553	1.772	0.0773	.
Carat:Depth	0.0095943	0.0060107	1.596	0.1114	

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2306 on 345 degrees of freedom

Multiple R-squared: 0.9302, Adjusted R-squared: 0.9292

F-statistic: 919.9 on 5 and 345 DF, p-value: < 2.2e-16

Call:

```
lm(formula = logPrice ~ Carat + Carat_sq, data = diamonds)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.8215	-0.1313	0.0003	0.1391	0.8615

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	6.13042	0.05218	117.48	<2e-16 ***
Carat	3.05963	0.08422	36.33	<2e-16 ***
Carat_sq	-0.52730	0.02944	-17.91	<2e-16 ***

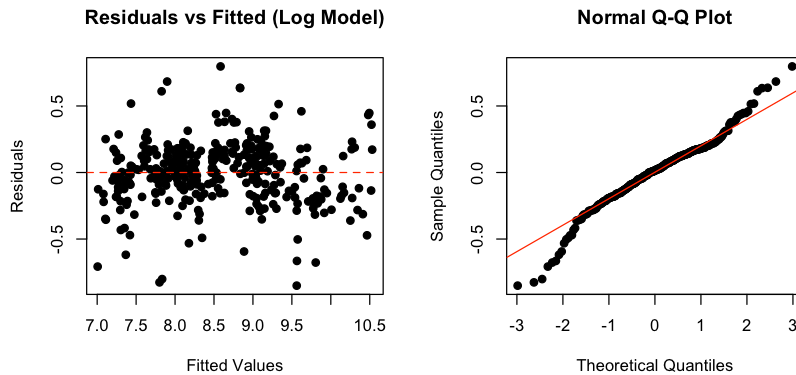
---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.238 on 348 degrees of freedom

Multiple R-squared: 0.925, Adjusted R-squared: 0.9246

F-statistic: 2146 on 2 and 348 DF, p-value: < 2.2e-16



Looking at the logPrice model, the model fit hasn't changed all that much, but the log transformation does seem to improve homoscedasticity and normality.

c. Once you have settled on a model for logPrice, produce similar graphs to those you found in (a). Has the log transformation helped with either the constant variance or normality conditions on the residuals?

Yes, the log transformation has helped with both, while retaining the fit of the model. Given the insignificance in interaction between carat and depth, I think we can fit a reduced model that leaves out the interaction while still factoring in the logPrice and squares of each of the carat and depth:

```
In [20]: log_reduced_model <- lm(logPrice ~ Carat + Depth + Carat_sq + Depth_sq,
                                data = diamonds)
summary(log_reduced_model)

par(mfrow=c(2,2))
plot(fitted(log_reduced_model), resid(log_reduced_model),
     xlab="Fitted Values",
     ylab="Residuals",
     main="Residuals vs Fitted (Reduced Log Model)",
     pch=19)
abline(h=0, col="red", lty=2)

qqnorm(resid(log_reduced_model), pch=19)
qqline(resid(log_reduced_model), col="red")
```

Call:

```
lm(formula = logPrice ~ Carat + Depth + Carat_sq + Depth_sq,
    data = diamonds)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.85847	-0.11996	0.01404	0.13568	0.80447

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	13.2641471	3.4062935	3.894	0.000118 ***
Carat	3.1151846	0.0827330	37.653	< 2e-16 ***
Depth	-0.2043145	0.1017786	-2.007	0.045480 *
Carat_sq	-0.5336499	0.0286216	-18.645	< 2e-16 ***
Depth_sq	0.0014319	0.0007547	1.897	0.058616 .

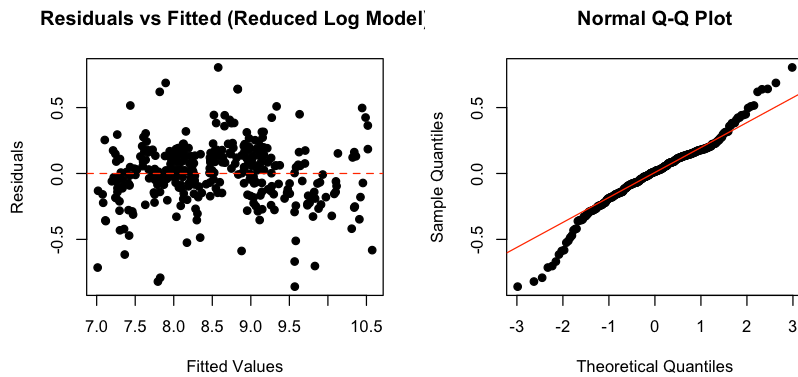
----

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2311 on 346 degrees of freedom

Multiple R-squared: 0.9297, Adjusted R-squared: 0.9289

F-statistic: 1144 on 4 and 346 DF, p-value: < 2.2e-16



## Problem 3.40

a. What average total price does the quadratic model predict for a .5-carat diamond?

```
In [22]: new_diamond <- data.frame(Carat = 0.5, Depth = mean(diamonds$Depth))
new_diamond$Carat_sq <- new_diamond$Carat^2

quad_model <- lm(TotalPrice ~ Carat + Carat_sq, data = diamonds)

predicted_price <- predict(quad_model, newdata = new_diamond)
print(predicted_price)
```

```
1
1794.843
```

b. Find a 95% confidence interval for the mean total price of .5-carat diamonds. Write a sentence interpreting the interval in terms that will make sense to the young couple.

```
In [23]: ci <- predict(quad_model, newdata = new_diamond, interval = "confidence", level = 0.95)
print(ci)
```

```
      fit      lwr      upr
1 1794.843 1424.296 2165.389
```

We can 95% confident that the average price for a .5 carat diamond falls between 1424 and 2165. Stated differently, while the price of a specific .5 carat diamond may cost more or less depending on their specific attributes, the true average price is very likely to fall between 1424 and 2165.



c. Find a 95% prediction interval for the total price when a diamond weighs .5 carat. Write a sentence interpreting the interval in terms that will make sense to the young couple.

```
In [24]: pred_int <- predict(quad_model, newdata = new_diamond, interval = "prediction")
print(pred_int)
```

```
      fit      lwr      upr
1 1794.843 -2404.462 5994.147
```

This output is a little bit tougher to interpret, since -2404 isn't meaningful in the real world (i.e., no seller would ever pay you

2400 to take the diamond), but puts simply, when shopping for a .5 – carat diamond, the price is likely to be between 0 to about \$6000. This wide range is a reflection of the fact that diamonds of the same carat can have their price impacted based on other factors (like cut, color and clarity, i.e., the other 3 Cs (I proposed to my now fiancée last year, and am all too familiar with these traits lol))

d. Repeat the previous two intervals (confidence and prediction) for the model found in part (b) of Exercise 3.39, where the response variable was logPrice. You should find the intervals for the log scale, but then exponentiate to give answers in terms of TotalPrice

```
In [26]: new_diamond <- data.frame(Carat = 0.5, Depth = mean(diamonds$Depth))
new_diamond$Carat_sq <- new_diamond$Carat^2
new_diamond$Depth_sq <- new_diamond$Depth^2

log_ci <- predict(log_reduced_model, newdata = new_diamond,
                  interval = "confidence", level = 0.95)
price_ci <- exp(log_ci)
print("95% Confidence Interval for mean price:")
print(price_ci)

log_pred <- predict(log_reduced_model, newdata = new_diamond,
                   interval = "prediction", level = 0.95)
price_pred <- exp(log_pred)
print("95% Prediction Interval for individual price:")
print(price_pred)
```

```
[1] "95% Confidence Interval for mean price:"
      fit      lwr      upr
1 1753.62 1656.519 1856.413
[1] "95% Prediction Interval for individual price:"
      fit      lwr      upr
1 1753.62 1109.174 2772.498
```

Using our reduced log model, we produce a 95% confidence interval of 1656 to 1856 to capture the mean price across all diamonds of .5 carats. We also produce 95% prediction interval of 1109 to 2772, which is a much more realistic range than our previous model (which included negative prices) of what a specific .5-carat diamond might actually cost.